

house-price-prediction

June 30, 2024

Boston House Price prediction

Importing Necessary Modelues

```
[1]: import pandas as pd
import numpy as np
import sklearn
from sklearn.impute import SimpleImputer
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

Importing Dataset for Analysis and Prediction

```
[2]: data = pd.read_csv('/kaggle/input/boston-house/HousingData.csv')
data.head()
```

```
[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

Describing Data and his Characterstics

```
[3]: data.columns
```

```
[3]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
        'PTRATIO', 'B', 'LSTAT', 'MEDV'],
```

```
dtype='object')
```

```
[4]: data.shape
```

```
[4]: (506, 14)
```

```
[5]: data.dtypes
```

```
[5]: CRIM      float64
     ZN       float64
     INDUS   float64
     CHAS     float64
     NOX      float64
     RM       float64
     AGE      float64
     DIS      float64
     RAD      int64
     TAX      int64
     PTRATIO  float64
     B        float64
     LSTAT    float64
     MEDV     float64
dtype: object
```

```
[6]: # Displaying Data Range
```

```
min_values = data.min()
max_values = data.max()
val = pd.DataFrame({'Min':min_values, 'Max':max_values})
print(val)
```

	Min	Max
CRIM	0.00632	88.9762
ZN	0.00000	100.0000
INDUS	0.46000	27.7400
CHAS	0.00000	1.0000
NOX	0.38500	0.8710
RM	3.56100	8.7800
AGE	2.90000	100.0000
DIS	1.12960	12.1265
RAD	1.00000	24.0000
TAX	187.00000	711.0000
PTRATIO	12.60000	22.0000
B	0.32000	396.9000
LSTAT	1.73000	37.9700
MEDV	5.00000	50.0000

Removing Nan Values

```
[7]: data.isna().sum()
```

```
[7]: CRIM      20  
     ZN        20  
     INDUS    20  
     CHAS      20  
     NOX       0  
     RM        0  
     AGE      20  
     DIS       0  
     RAD       0  
     TAX       0  
     PTRATIO   0  
     B         0  
     LSTAT     20  
     MEDV      0  
     dtype: int64
```

```
[8]: imp_mean = SimpleImputer(strategy='mean')  
     imp_median = SimpleImputer(strategy='median')  
     imp_mode = SimpleImputer(strategy='most_frequent')
```

```
[9]: data['CRIM'] = imp_median.fit_transform(data[['CRIM']])
```

```
[10]: data['ZN'] = imp_mode.fit_transform(data[['ZN']])
```

```
[11]: data['INDUS'] = imp_mean.fit_transform(data[['INDUS']])
```

```
[12]: data['CHAS'] = imp_mode.fit_transform(data[['CHAS']])
```

```
[13]: data['AGE'] = imp_median.fit_transform(data[['AGE']])
```

```
[14]: data['LSTAT'] = imp_mean.fit_transform(data[['LSTAT']])
```

```
[15]: data.isna().sum()
```

```
[15]: CRIM      0  
     ZN        0  
     INDUS    0  
     CHAS      0  
     NOX       0  
     RM        0  
     AGE      0  
     DIS       0  
     RAD       0  
     TAX       0  
     PTRATIO   0
```

```

B          0
LSTAT      0
MEDV       0
dtype: int64

```

```

[16]: # Data after Preprocessing
data.head()

```

```

[16]:      CRIM      ZN  INDUS  CHAS    NOX     RM   AGE     DIS  RAD  TAX  PTRATIO  \
0  0.00632  18.0    2.31   0.0  0.538  6.575  65.2  4.0900    1  296     15.3
1  0.02731   0.0    7.07   0.0  0.469  6.421  78.9  4.9671    2  242     17.8
2  0.02729   0.0    7.07   0.0  0.469  7.185  61.1  4.9671    2  242     17.8
3  0.03237   0.0    2.18   0.0  0.458  6.998  45.8  6.0622    3  222     18.7
4  0.06905   0.0    2.18   0.0  0.458  7.147  54.2  6.0622    3  222     18.7

      B      LSTAT  MEDV
0  396.90  4.980000  24.0
1  396.90  9.140000  21.6
2  392.83  4.030000  34.7
3  394.63  2.940000  33.4
4  396.90 12.715432  36.2

```


Visualizing to describe the data

```

[17]: # Box plot of CRIM


plt.figure(figsize=(5,4))
sns.boxplot(x=data['CRIM'])
plt.title('Box Plot of CRIM')
plt.xlabel('Per Capita Crime Rate by Town (CRIM)')
plt.show()

```



```
[18]: # Bar plot of CHAS

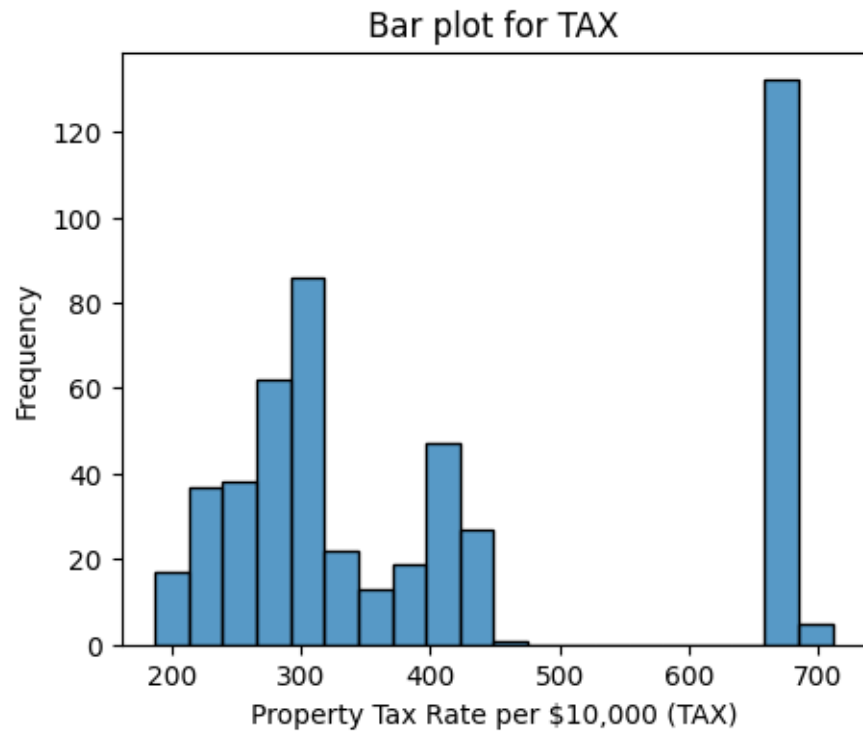
plt.figure(figsize=(5,4))
sns.countplot(x='CHAS', data=data)
plt.title('Count of CHAS')
plt.xlabel('Charles River Dummy Variable (CHAS)')
plt.ylabel('Count')
plt.show()
```



```
[19]: # Bar Plot for Tax


plt.figure(figsize=(5, 4))
sns.histplot(data['TAX'], bins=20)
plt.title('Bar plot for TAX')
plt.xlabel('Property Tax Rate per $10,000 (TAX)')
plt.ylabel('Frequency')
plt.show()
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```



```
[20]: #Correlation matrix

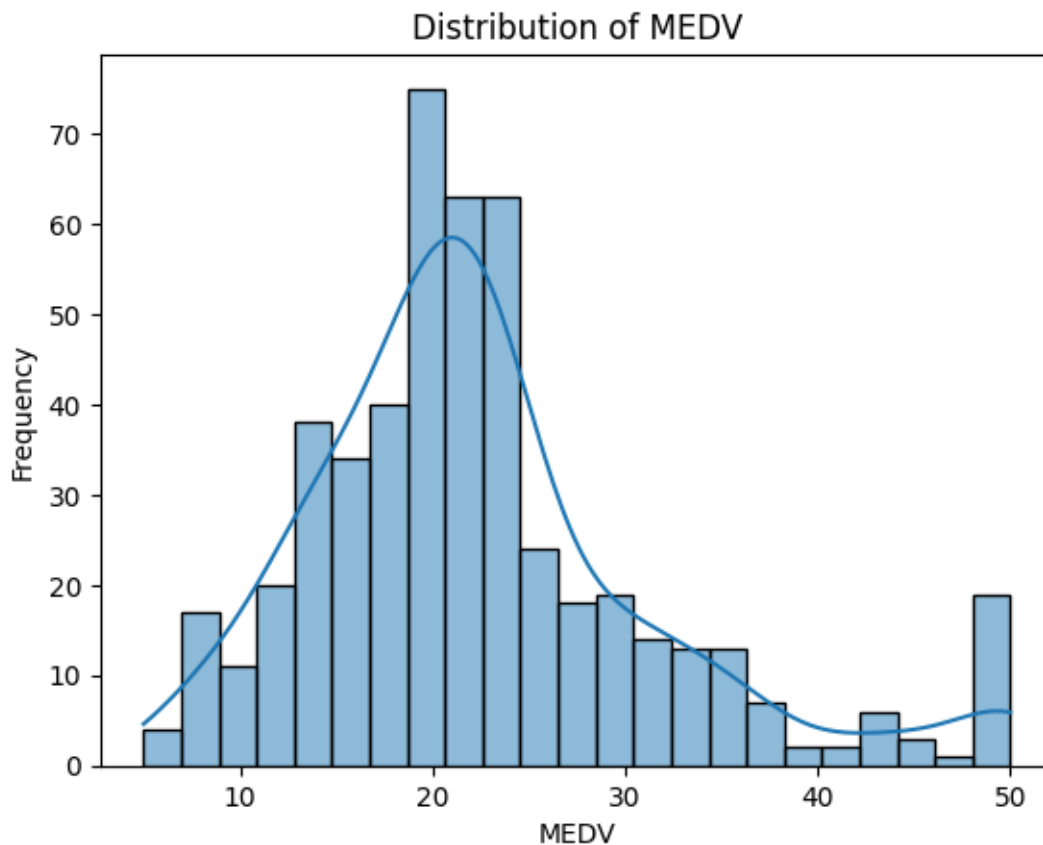
correlation_matrix = data.corr().round(2)
sns.heatmap(data=correlation_matrix, annot=True, cmap='coolwarm')
plt.show()
```



[21]: *# Histogram for Medv*


```
sns.histplot(data['MEDV'], kde=True)
plt.xlabel('MEDV')
plt.ylabel('Frequency')
plt.title('Distribution of MEDV')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



```
[22]: # Violin plot for Medv and CHAS

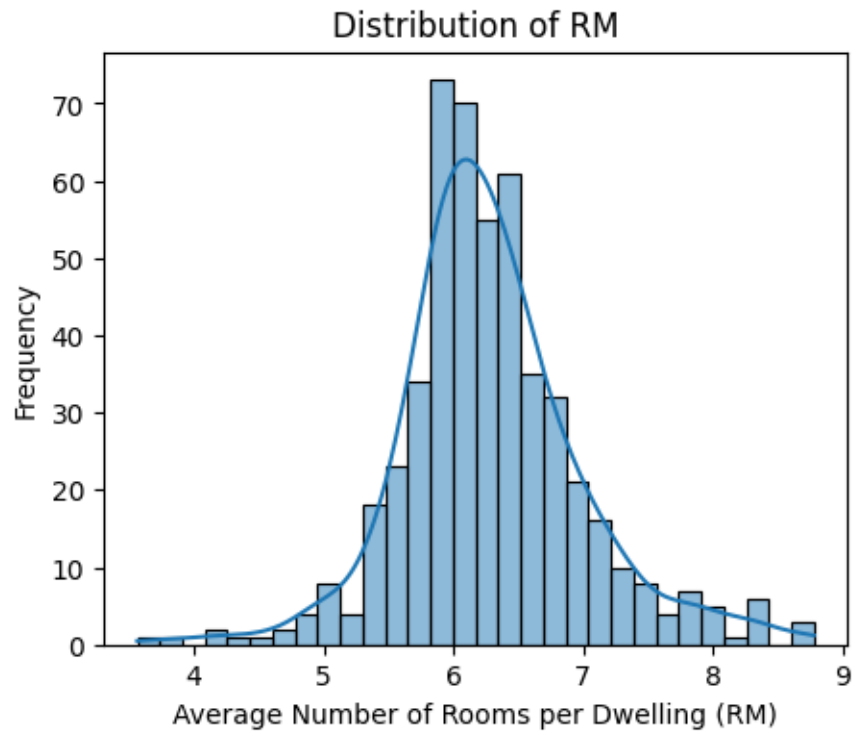
plt.figure(figsize=(5,4))
sns.violinplot(x='CHAS', y='MEDV', data=data)
plt.title('MEDV Distribution by CHAS')
plt.xlabel('Charles River Dummy Variable (CHAS)')
plt.ylabel('Median Value of Homes (MEDV)')
plt.show()
```



```
[23]: # Histogram of RM



plt.figure(figsize=(5,4))
sns.histplot(data['RM'], bins=30, kde=True)
plt.title('Distribution of RM')
plt.xlabel('Average Number of Rooms per Dwelling (RM)')
plt.ylabel('Frequency')
plt.show()
```

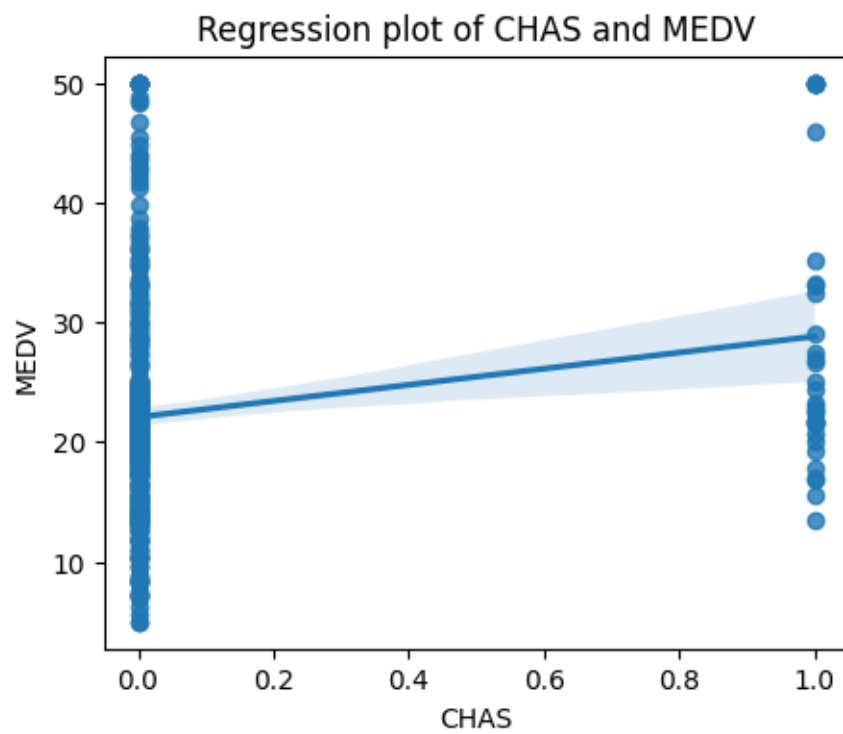

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```





[24]: *# Relationship between Average house price and other Features*

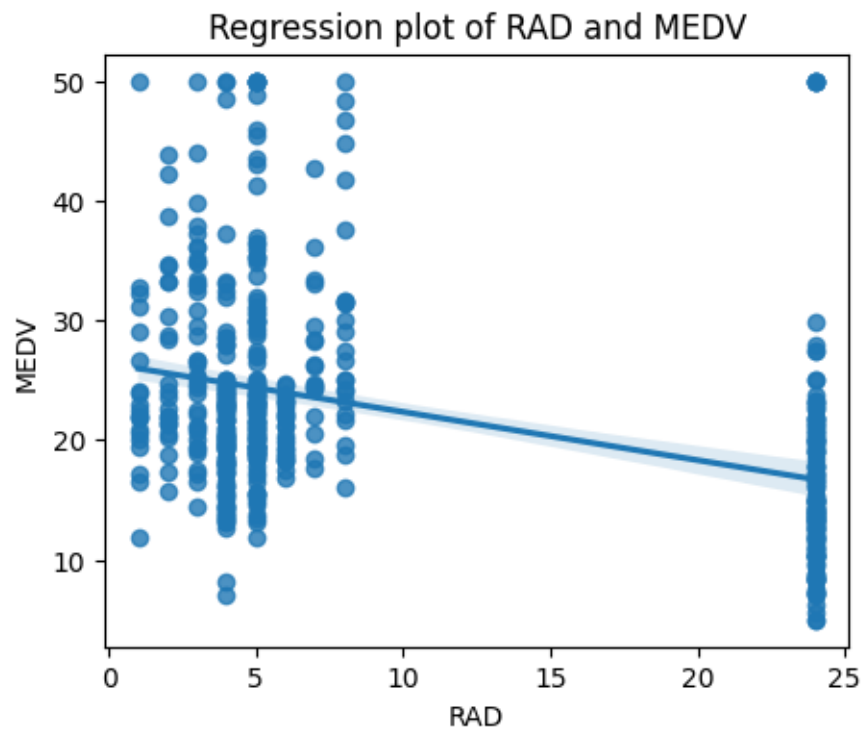
```
for column in data.columns[:-1]:  
    plt.figure(figsize=(5,4))  
    sns.regplot(x=data[column], y=data['MEDV'])  
    plt.title(f'Regression plot of {column} and MEDV')  
    plt.xlabel(column)  
    plt.ylabel('MEDV')  
    plt.show()
```












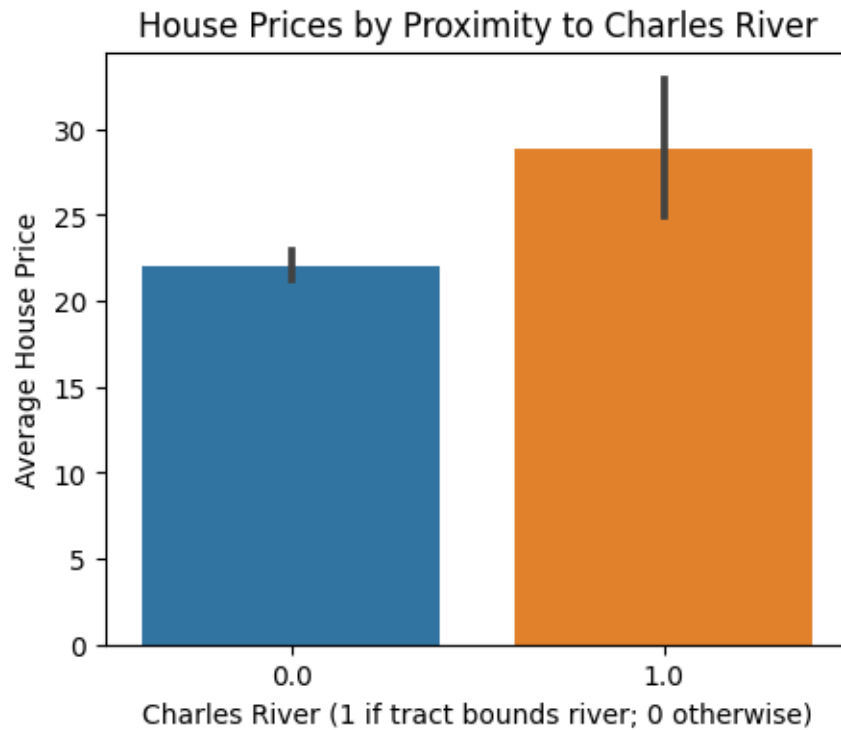







```
[25]: # Average house prices by proximity to the Charles River

plt.figure(figsize=(5,4))
sns.barplot(x='CHAS', y='MEDV', data=data)
plt.title('House Prices by Proximity to Charles River')
plt.xlabel('Charles River (1 if tract bounds river; 0 otherwise)')
plt.ylabel('Average House Price')
plt.show()
```



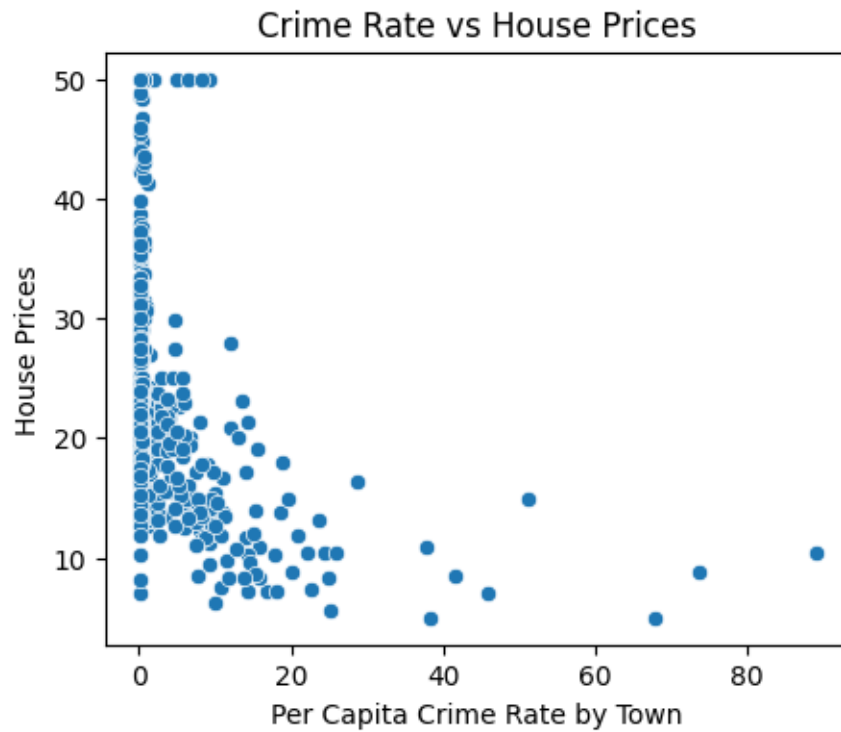
[26]: *# Relationship between RM (average number of rooms per dwelling) and MEDV ↴
↴(house price)*

```
plt.figure(figsize=(5,4))
sns.scatterplot(x='RM', y='MEDV',data=data)
plt.title('Number of Rooms vs House Prices')
plt.xlabel('Average Number of Rooms per Dwelling')
plt.ylabel('House Prices')
plt.show()
```




```
[27]: # Relationship between CRIM (per capita crime rate by town) and MEDV (house
      ↪ price)
```

```
plt.figure(figsize=(5,4))
sns.scatterplot(x='CRIM', y='MEDV', data=data)
plt.title('Crime Rate vs House Prices')
plt.xlabel('Per Capita Crime Rate by Town')
plt.ylabel('House Prices')
plt.show()
```




[28]: *# Relationship between PTRATIO (pupil-teacher ratio by town) and MEDV (house*
↪price)

```
plt.figure(figsize=(5,4))
sns.scatterplot(x='PTRATIO', y='MEDV', data=data)
plt.title('Pupil-Teacher Ratio vs House Prices')
plt.xlabel('Pupil-Teacher Ratio by Town')
plt.ylabel('House Prices')
plt.show()
```



```
[29]: # Relationship between AGE and AVERAGE house price

plt.figure(figsize=(5,4))
sns.scatterplot(x='AGE', y='MEDV', data=data)
plt.title('Age of House vs House Prices')
plt.xlabel('Proportion of Owner-Occupied Units Built Prior to 1940')
plt.ylabel('House Prices')
plt.show()
```



Building Machine Learning Models

```
[30]: # Splitting data for Training and Testing

x = data.drop('MEDV', axis=1)
y = data['MEDV']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=42)
```

```
[31]: # Simple Linear Regression Model

model = LinearRegression()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
```

```
[32]: # Model Evaluationing

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```


```
[33]: print(f'Mean Absolute Error (MAE): {mae}')
      print(f'Mean Squared Error (MSE): {mse}')
      print(f'Root Mean Squared Error (RMSE): {rmse}')
      print(f'R-squared (R2): {r2}')
```

Mean Absolute Error (MAE): 3.1584994146197096
Mean Squared Error (MSE): 25.072290196306753
Root Mean Squared Error (RMSE): 5.007223801300153
R-squared (R2): 0.6581072308584777

```
[34]: #Evaluation metrics for Linear Regression Model

metrics = {'MAE': mae, 'MSE': mse, 'RMSE': rmse, 'R-squared (R2) ': r2}
metric_names = list(metrics.keys())
metric_values = list(metrics.values())
plt.figure(figsize=(5,4))
sns.barplot(x=metric_names, y=metric_values, palette='viridis')
plt.title('Evaluation Metrics for Linear Regression Model')
plt.xlabel('Metrics')
plt.ylabel('Values')
plt.show()
```

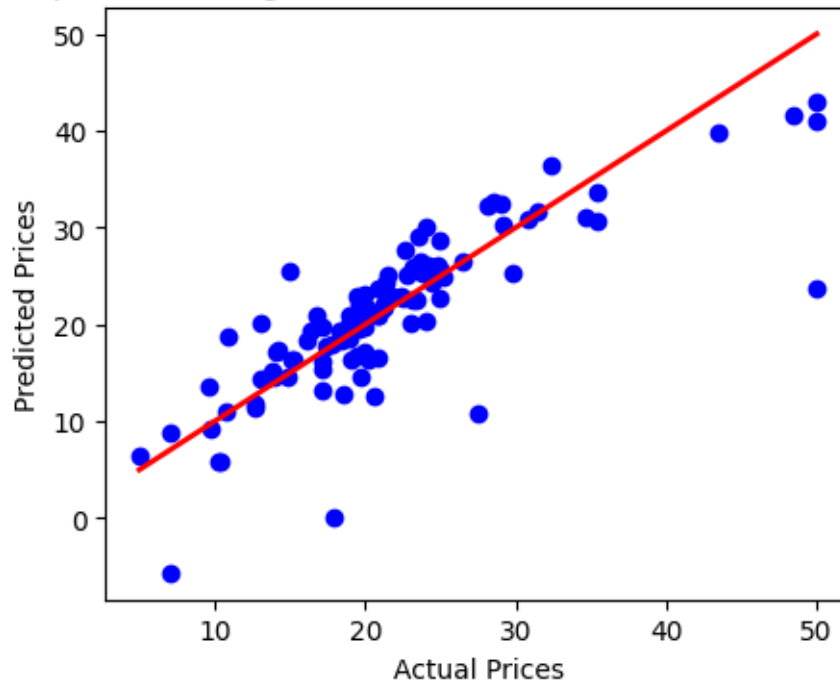
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1765: FutureWarning:
unique with argument that is not not a Series, Index, ExtensionArray, or
np.ndarray is deprecated and will raise in a future version.
order = pd.unique(vector)




```
[35]: # Multiple Linear Regression

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         ↪linewidth=2)
plt.title('Multiple Linear Regression for Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```

Multiple Linear Regression for Predicted vs Actual House Prices




```
[36]: #Scaling the features

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         ↪linewidth=2)
plt.title('Multiple Linear Regression with Scaled Features: Predicted vs Actual_
         ↪House Prices')
plt.xlabel('Actual Prices')
```

```
plt.ylabel('Predicted Prices')
plt.show()
```

Multiple Linear Regression with Scaled Features: Predicted vs Actual House Prices



```
[37]: #predict house prices using a Decision Tree Regressor

from sklearn.tree import DecisionTreeRegressor

tree_model = DecisionTreeRegressor(random_state=42)
tree_model.fit(x_train, y_train)
y_pred = tree_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R-squared (R2) Score:", r2_score(y_test, y_pred))
```


Mean Squared Error: 12.571274509803924
R-squared (R2) Score: 0.8285745809360402

```
[38]: # Decision Tree Regression for predicted and Actual Prices

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         ↪linewidth=2)
plt.title('Decision Tree Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
```

```
plt.show()
```

Decision Tree Regression: Predicted vs Actual House Prices



```
[39]: #predict house prices using a Random Forest Regressor

from sklearn.ensemble import RandomForestRegressor

forest_model = RandomForestRegressor(random_state=42, n_estimators=100)
forest_model.fit(x_train, y_train)
y_pred = forest_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R-squared (R2) Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 8.278940480392155


R-squared (R2) Score: 0.8871060495775506

```
[40]: # Random Forest Regression for Predicted and Actual prices

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Random Forest Regression: Predicted vs Actual House Prices')
```

```
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```

Random Forest Regression: Predicted vs Actual House Prices



```
[41]: # predict house prices using a Gradient Boosting Regressor

from sklearn.ensemble import GradientBoostingRegressor

gb_model = GradientBoostingRegressor(random_state=42, n_estimators=100)
gb_model.fit(x_train, y_train)
y_pred = gb_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R-squared (R2) Score:", r2_score(y_test, y_pred))
```


Mean Squared Error: 7.155319376730123
R-squared (R2) Score: 0.9024280615512894

```
[42]: # Gradient Boosting Regression for Predicted and Actual prices

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
```

```
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Gradient Boosting Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```

Gradient Boosting Regression: Predicted vs Actual House Prices



[43]: *#predict house prices using a Support Vector Regressor*

```
from sklearn.svm import SVR

svr_model = SVR(kernel='rbf')
svr_model.fit(x_train, y_train)
y_pred = svr_model.predict(x_test)


print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R-squared (R2) Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 52.931033106567114
R-squared (R2) Score: 0.2782176123267932

[44]: *# Support Vector Regression for Predicted and Actual prices*

```
plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Support Vector Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```

Support Vector Regression: Predicted vs Actual House Prices



[45]: *#predict house prices using a K-Nearest Neighbors Regressor*

```
from sklearn.neighbors import KNeighborsRegressor

knn_model = KNeighborsRegressor(n_neighbors=5)
knn_model.fit(x_train, y_train)
y_pred = knn_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R-squared (R2) Score:", r2_score(y_test, y_pred))
```


Mean Squared Error: 29.704478431372554

R-squared (R2) Score: 0.5949414151124269

```
[46]: # K-Nearest Neighbors Regression for Predicted and Actual prices

plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('K-Nearest Neighbors Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```

K-Nearest Neighbors Regression: Predicted vs Actual House Prices



```
[47]: # predict house prices using Ridge Regression
from sklearn.linear_model import Ridge


ridge_model = Ridge(alpha=1.0)
ridge_model.fit(x_train, y_train)
y_pred = ridge_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 25.2687878117523
R² Score: 0.6554277343566272

```
[48]: # Ridge Regression for Predicted and Actual prices

plt.figure(figsize=(5, 4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Ridge Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```



```
[49]: # predict house prices using Lasso Regression

from sklearn.linear_model import Lasso

lasso_model = Lasso(alpha=0.1)
lasso_model.fit(x_train, y_train)
y_pred = lasso_model.predict(x_test)


print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 25.72864216726899

R² Score: 0.6491570316095543

```
[50]: # Lasso Regression for Predicted and Actual prices

plt.figure(figsize=(5, 4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Lasso Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```



```
[51]: # predict house prices using Elastic Net Regression

from sklearn.linear_model import ElasticNet


elastic_model = ElasticNet(alpha=0.1, l1_ratio=0.5)
elastic_model.fit(x_train, y_train)
y_pred = elastic_model.predict(x_test)

print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 24.91093992270312
R² Score: 0.6603074483660197

```
[52]: # Elastic Net Regression for Predicted and Actual prices

plt.figure(figsize=(5, 4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('Elastic Net regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```



```
[53]: # predict house prices using XGBoost Regressor

import xgboost as xgb

xgb_model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=100,
                             random_state=42)
xgb_model.fit(x_train, y_train)
y_pred = xgb_model.predict(x_test)
```


```
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
```

Mean Squared Error: 5.537024472117028

R² Score: 0.9244955839791882

```
[54]: # XGBoost Regression for Predicted and Actual prices
```

```
plt.figure(figsize=(5,4))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linewidth=2)
plt.title('XGBoost Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()
```



```
[55]: # PCA for Dimensionality Reduction
```

```
from sklearn.decomposition import PCA

pca = PCA(n_components=10)
x_pca = pca.fit_transform(x)
```

```

xp_train, xp_test, yp_train, yp_test = train_test_split(x_pca, y, test_size=0.
↳2, random_state=42)

linear_model = LinearRegression()
linear_model.fit(xp_train, yp_train)
yp_pred = linear_model.predict(xp_test)

print("Mean Squared Error:", mean_squared_error(yp_test, yp_pred))
print("R^2 Score:", r2_score(yp_test, yp_pred))

```

Mean Squared Error: 24.264297405017974
R² Score: 0.6691252467914937


[56]: *# PCA for Predicted and Actual prices*

```

plt.figure(figsize=(5,4))
plt.scatter(yp_test, yp_pred, color='blue')
plt.plot([min(yp_test), max(yp_test)], [min(yp_test), max(yp_test)],
↳color='red', linewidth=2)
plt.title('PCA with Linear Regression: Predicted vs Actual House Prices')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.show()

```

PCA with Linear Regression: Predicted vs Actual House Prices



```
[57]: # compare the performance of different models

from sklearn.model_selection import cross_val_score

models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(random_state=42),
    'Random Forest': RandomForestRegressor(random_state=42, n_estimators=100),
    'Gradient Boosting': GradientBoostingRegressor(random_state=42,
    ↪n_estimators=100),
    'Support Vector Regressor': SVR(kernel='rbf')
}

# Compare models using cross-validation
results = {}
for model_name, model in models.items():
    cv_scores = cross_val_score(model, x, y, cv=5, scoring='r2')
    results[model_name] = cv_scores

# Print the results
for model_name, cv_scores in results.items():
    print(f"{model_name}: Mean R-squared (R2) Score = {np.mean(cv_scores):.3f},
    ↪Std = {np.std(cv_scores):.3f}")

# Dictionary to store evaluation metrics
evaluation_metrics = {
    'Model': [],
    'MAE': [],
    'MSE': [],
    'RMSE': [],
    'R-Squared': []
}

# Train and evaluate each model
for name, model in models.items():
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)

    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    evaluation_metrics['Model'].append(name)
    evaluation_metrics['MAE'].append(mae)
```

```

evaluation_metrics['MSE'].append(mse)
evaluation_metrics['RMSE'].append(rmse)
evaluation_metrics['R-Squared'].append(r2)

```

Linear Regression: Mean R-squared (R2) Score = 0.340, Std = 0.392
 Decision Tree: Mean R-squared (R2) Score = 0.221, Std = 0.476
 Random Forest: Mean R-squared (R2) Score = 0.633, Std = 0.177
 Gradient Boosting: Mean R-squared (R2) Score = 0.675, Std = 0.112
 Support Vector Regressor: Mean R-squared (R2) Score = -0.093, Std = 0.283

```

[58]: # Convert the evaluation metrics dictionary to a DataFrame
evaluation_df = pd.DataFrame(evaluation_metrics)
evaluation_df

```

```

[58]:

```

	Model	MAE	MSE	RMSE	R-Squared
0	Linear Regression	3.158499	25.072290	5.007224	0.658107
1	Decision Tree	2.732353	12.571275	3.545599	0.828575
2	Random Forest	2.050010	8.278940	2.877315	0.887106
3	Gradient Boosting	1.899569	7.155319	2.674943	0.902428
4	Support Vector Regressor	4.524912	52.931033	7.275372	0.278218

```

[59]: # Plot all evaluation metrics
fig, axs = plt.subplots(2, 2, figsize=(16,10))

# MAE plot
sns.barplot(x='Model', y='MAE', data=evaluation_df, palette='viridis',
            ax=axs[0, 0])
axs[0, 0].set_title('Mean Absolute Error (MAE) of Models')
axs[0, 0].set_xlabel('Model')
axs[0, 0].set_ylabel('MAE')

# MSE plot
sns.barplot(x='Model', y='MSE', data=evaluation_df, palette='viridis',
            ax=axs[0, 1])
axs[0, 1].set_title('Mean Squared Error (MSE) of Models')
axs[0, 1].set_xlabel('Model')
axs[0, 1].set_ylabel('MSE')

# RMSE plot
sns.barplot(x='Model', y='RMSE', data=evaluation_df, palette='viridis',
            ax=axs[1, 0])
axs[1, 0].set_title('Root Mean Squared Error (RMSE) of Models')
axs[1, 0].set_xlabel('Model')
axs[1, 0].set_ylabel('RMSE')

# R² plot


```

```

sns.barplot(x='Model', y='R-Squared', data=evaluation_df, palette='viridis',
            ax=axes[1, 1])
axes[1, 1].set_title('R-squared of Models')
axes[1, 1].set_xlabel('Model')
axes[1, 1].set_ylabel('R2')

plt.tight_layout()
plt.show()

```



[60]: *# Clustering the data using KMeans Clustering*

```

from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=3, random_state=42)
data['Cluster'] = kmeans.fit_predict(x)

```

/opt/conda/lib/python3.10/site-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

[61]: *# Visualize the Clustered Data using K-Means Clustering*

```

custom_palette = ['#000000', '#00FF00', '#0000FF']
plt.figure(figsize=(5,4))

```

```

sns.scatterplot(x='RM', y='MEDV', hue='Cluster', palette=custom_palette, data=data)
plt.title('KMeans Clustering: Neighborhoods')
plt.xlabel('Average Number of Rooms per Dwelling')
plt.ylabel('House Prices')
plt.show()

```

